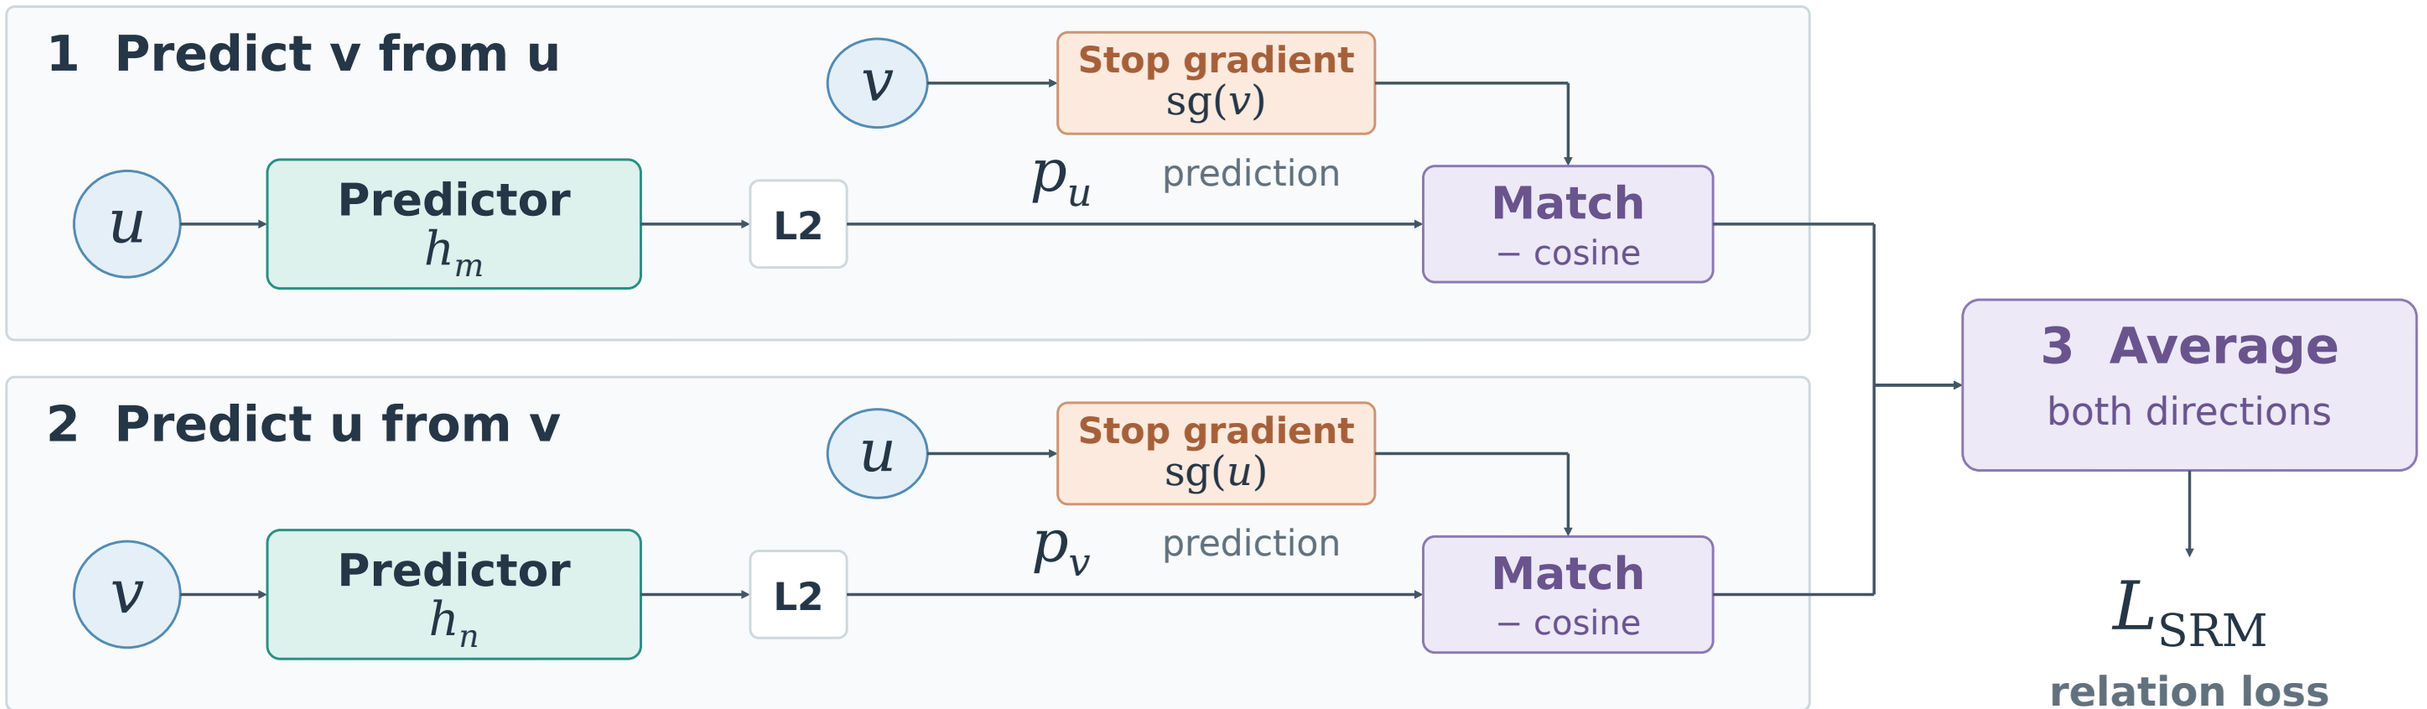
SRM | Symmetric feature matching

Predict the other feature → swap roles → average the two losses.



Same inputs, opposite roles. Only the prediction path receives gradients.

USE THE SAME MODULE FOR FOUR RELATIONS

Image ↔ Image		Image ↔ Prompt		Image ↔ Answer		Prompt ↔ Answer	
Input pair	Predictors	Input pair	Predictors	Input pair	Predictors	Input pair	Predictors
(v_1,v_2)	$(h_{\mathrm{I}},h_{\mathrm{I}})$	(a^c,t_P)	$(h_{\mathrm{I}},h_{\mathrm{T}})$	(a^c,t_A)	$(h_{\mathrm{I}},h_{\mathrm{T}})$	(t_P,t_A)	$(h_{\mathrm{T}},h_{\mathrm{T}})$

Only two predictor networks: one for images and one for text, reused across all four relations.

Current mode: positive matching only · 256-D inputs · PI/PT are outside SRM · Targets are detached, not permanently frozen