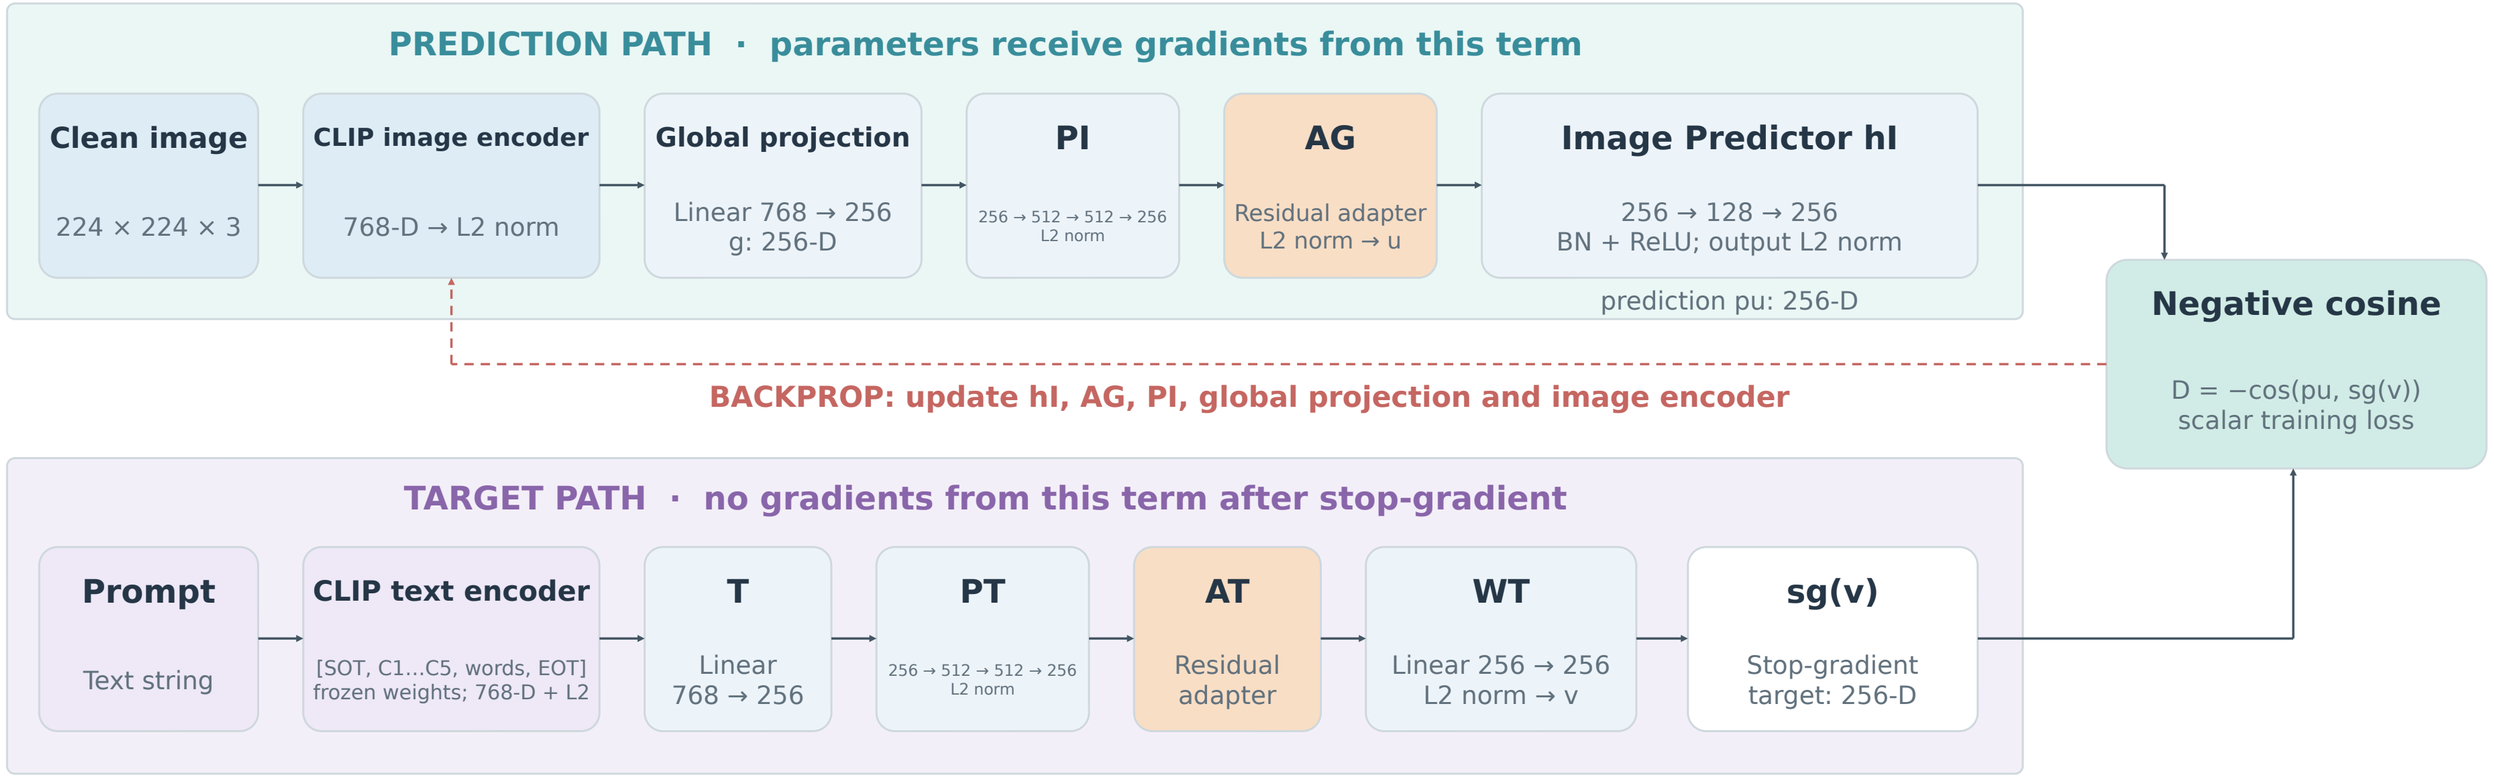
SRM: Image predicts Prompt

One direction of the image–prompt pair · projection_only_v1 baseline



sg blocks this loss from updating C1…C5, T, PT, AT and WT on the target path.

Those text parameters can still learn from the reverse direction (Prompt predicts Image) and other losses.
CLIP text weights stay frozen in every direction; five context vectors are shared by Prompt and Answer.

Solid arrows: forward flow Red dashed arrow: gradient direction L2 norm: vector normalization No local P0, QA or final scoring head in this path.